

# Statistics 452: Statistical Learning and Prediction

## Chapter 8, Part 3: Boosting

Brad McNeney

2017-11-06

# Introduction to Boosting

- ▶ Reference: Hastie, Tibshirani and Friedman (2001). The Elements of Statistical Learning (hereafter ESL).
- ▶ Motivation for boosting: Combine many “weak” classifiers to produce a powerful “committee”.
  - ▶ Similar in this respect to bagging, but otherwise fundamentally different.
- ▶ A weak classifier is one that does little better than guessing.
  - ▶ On its own a weak classifier is not useful, but if applied *sequentially*, it can produce a powerful classifier.

## Example Boosting Algorithm: AdaBoost.M1

- ▶ Due to Freund and Schapire (1997).
- ▶ Suppose two outcome classes  $Y = -1$  or  $1$  and a “base” classifier that produces a prediction.
  - ▶ Need not be a decision tree classifier at this point.
- ▶ Sequentially apply the classifier to modified versions of the data (more on next slide), leading to a sequence of weak classifiers  $G_m(x)$ ;  $m = 1, \dots, M$  which are weighted to give final predictions.

# AdaBoost Weighting

- ▶ Combine predictions with a weighted majority vote

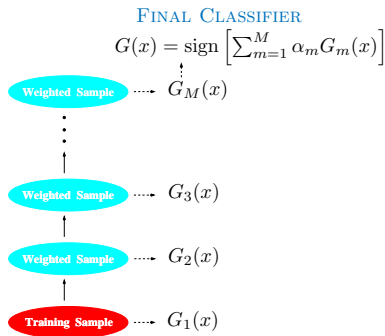
$$G(x) = \text{sign} \left( \sum_{m=1}^M \alpha_m G_m(x) \right),$$

which classifies as 1 if weighted sum  $> 0$  and  $-1$  otherwise.

- ▶ The classifier weights  $\alpha_m$  are computed by the algorithm to give higher weight to more accurate classifiers.
- ▶ Modify the data at each boosting step by applying observation weights  $w_1, \dots, w_n$ .
  - ▶ Initially all weights are equal.
  - ▶ At step  $m$ , observations that were misclassified at step  $m - 1$  are up-weighted.
  - ▶ As we go, observations that are difficult to classify receive more and more weight, forcing the weak classifier to focus on them.
- ▶ Full details in Algorithm 10.1 of ESL (page 301).

# Schematic

Elements of Statistical Learning (2nd Ed.) ©Hastie, Tibshirani & Friedman 2009 Chap 10



**FIGURE 10.1.** *Schematic of AdaBoost. Classifiers are trained on weighted versions of the dataset, and then combined to produce a final prediction.*

# AdaBoost as an Additive Model

- ▶ Let  $b(x; \gamma)$  be the base classifier for parameters  $\gamma$ .
  - ▶ Let  $\gamma_m$  denote the values at step  $m$ , so that  $G_m(x) = b(x; \gamma_m)$  is the classifier at step  $m$ . This is a basis function.
- ▶ The classifier weights are the coefficients of the basis functions.
- ▶ The additive model is

$$f(x; \alpha, \gamma) = \sum_{m=1}^M \alpha_m b(x; \gamma_m)$$

- ▶ We would like to find the coefficients  $\alpha = (\alpha_1, \dots, \alpha_M)$  and  $\gamma = (\gamma_1, \dots, \gamma_M)$  that minimize a “loss function”,

$$\sum_{i=1}^n L(y_i, f(x_i; \alpha)).$$

- ▶ We are used to the squared-error loss  $L(y, f(x)) = (y - f(x))^2$ , but others are possible.

# Forward Stagewise Additive Modelling

- ▶ Approximate the solution by a greedy algorithm that sequentially adds the “best” new basis function, without adjusting the coefficients of those previously added.
  1. Initialize  $f_0(x) = 0$ .
  2. For  $m = 1 : M$ 
    - 2.1 Find the  $\alpha_m$  and  $\gamma_m$  that minimize 
$$\sum_{i=1}^n L(y_i, f_{m-1}(x_i) + \alpha b(x_i; \gamma))$$
    - 2.2 Set  $f_m(x) = f_{m-1}(x) + \alpha_m b(x; \gamma_m)$
  3. Return  $\hat{f}(x) = f_M(x)$ .

## Example Forward Stagewise Additive Model

- ▶ One can show (ESL Section 10.4) that AdaBoost is forward stagewise additive modelling with the exponential loss function  $L(y, f(x)) = \exp(-yf(x))$ .



# Boosting Decision Trees

- ▶ The parameters of a decision tree are the disjoint regions (obtained by recursive partitioning) and the values assigned to each region.
- ▶ Let  $T(x; \gamma)$  be a tree.
- ▶ The boosted tree model is a sum

$$f_M(x) = \sum_{m=1}^M T(x; \gamma_m)$$

(no weighting), where the trees at step  $m$  are fit according to the forward stagewise algorithm.

- ▶ At step  $m$  we find the  $\gamma_m$  that minimizes

$$\sum_{i=1}^n L(y_i, f_{m-1}(x_i) + T(x_i; \gamma)) \quad (1)$$

and take  $f_m(x) = f_{m-1}(x) + T(x; \gamma_m)$ .

# Boosting Regression Trees

- ▶ If a regression tree and the loss is squared-error loss,

$$\begin{aligned}L(y_i, f_{m-1}(x_i) + T(x_i; \gamma)) &= (y_i - f_{m-1}(x_i) - T(x_i; \gamma))^2 \\ &= (r_i^{(m-1)} - T(x_i; \gamma))^2,\end{aligned}$$

where  $r_i^{(m-1)}$  is the  $i$ th residual from step  $m - 1$ .

- ▶ Solve (1) by fitting a tree to the residuals (Our text, Alg. 8.2).
- ▶ Note: As a basis function,  $T(x; \gamma)$  could, in general, depend on all predictors, which would make the boosted model not additive in the sense of Chapter 7.
  - ▶ When the trees have only two leaves (i.e., one split on one variable), the boosted model is additive in the sense of Chapter 7.

# Gradient Boosting

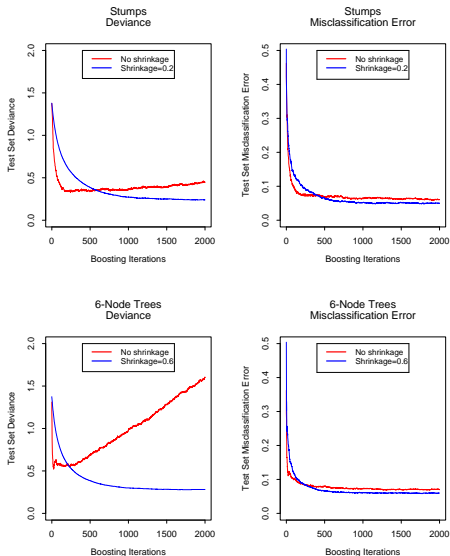
- ▶ With loss functions other than squared-error and exponential, the solution to (1) is more challenging.
- ▶ A general, but approximate algorithm based on ideas from optimization is called gradient boosting.
  - ▶ A description is beyond the scope of this course.
  - ▶ We use the implementation in the `gbm` package.
  - ▶ A more modern implementation of boosting trees is in the package `xgboost` (eXtreme Gradient Boosting). More flexible, computations highly optimized and parallelized, but same basic idea.

# Choosing the Depth of the Trees

- ▶ Set the tree depth to be the same for all trees.
- ▶ Could consider the depth as a tuning parameter and choose it by cross-validation.
- ▶ Text and software suggest  $d = 1$  is often fine.
  - ▶ Software calls  $d$  the interaction depth. For  $d > 1$  each tree depends on more than one variable and would represent an “interaction”.

# Shrinkage

- ▶ Large  $M$  will lead to overfitting.
- ▶ Can select  $M$  as a tuning parameter, but experience has shown that it is better to take a large  $M$  and shrink the contributions of each tree by a factor  $\lambda$ ; that is, take
$$f_m(x) = f_{m-1}(x) + \lambda T(x; \gamma_m).$$

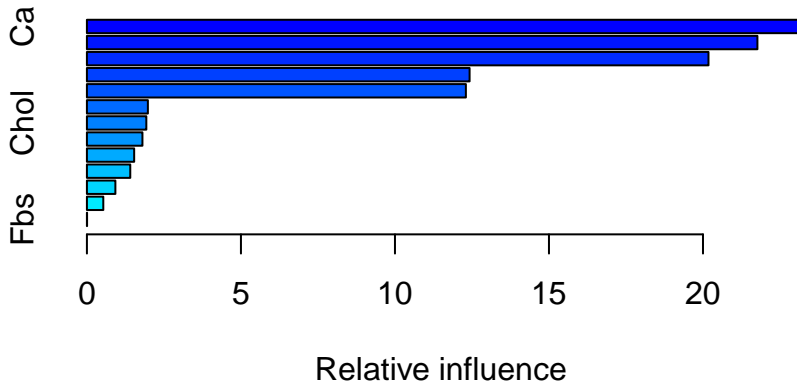


**FIGURE 10.11.** Test error curves for simulated example (10.2) of Figure 10.9, using gradient boosting (MART). The models were trained using binomial deviance, either stumps or six terminal-node trees, and with or without shrinkage. The left panels report test

## Example: Heart Data

- ▶ Recall that the best tree fit the to Heart data had test-set misclassification rate about 27%,
- ▶ Random forest had a test-set misclassification of about 17%.

```
library(gbm)
hboost <- gbm(I(AHD=="Yes") ~ ., data=Heart[train,],
              n.trees=5000,distributio="bernoulli")
summary(hboost)
```



```
##          var      rel.inf
## Ca          Ca 23.215816764
## Thal        Thal 21.766668524
## ChestPain ChestPain 20.178691076
## Oldpeak     Oldpeak 12.423059066
## MaxHR       MaxHR 12.302340643
## Age         Age  1.078875272
```



```
boo.hpred <- predict(hboost,newdata=Heart[-train,],
                    n.trees=5000,type="response")
boo.hpred <- (boo.hpred>0.5)
table(boo.hpred,Heart[-train,]$AHD)
```

```
##
## boo.hpred No Yes
##      FALSE 50  12
##      TRUE  4  33
```

```
16/nrow(Heart[-train,]) # Lowest so far.
```

```
## [1] 0.1616162
```

# Change Shrinkage

```
hboost <- gbm(I(AHD=="Yes") ~ ., data=Heart[train,],  
              n.trees=5000,distribution="bernoulli",shrinkage=.2)  
boo.hpred <- predict(hboost,newdata=Heart[-train,],  
                     n.trees=5000,type="response")  
boo.hpred <- (boo.hpred>0.5)  
table(boo.hpred,Heart[-train,]$AHD)
```

```
##  
## boo.hpred No Yes  
##      FALSE 47  13  
##      TRUE  7  32
```

```
16/nrow(Heart[-train,]) # Worse
```

```
## [1] 0.1616162
```